

Praxish: A Rational Reconstruction of a Logic-Based DSL for Modeling Social Practices

James Dameris*, Rosaura Hernandez Roman*, Max Kreminski

Santa Clara University
500 El Camino Real
Santa Clara, California 95053 USA
{jdameris, rhernandez2, mkreminski}@scu.edu

Abstract

The Versu framework is historically notable for its full-featuredness as a suite of tools for creating highly responsive interactive dramas. However, it has also been lost for nearly a decade, and a similarly approachable and flexible simulationist interactive narrative authoring framework has not yet emerged to take its place. We therefore aim to introduce an open-source rational reconstruction of the Versu framework, drawing on publicly available documentation of Versu’s design and implementation to assemble a successor system with similar architecture and capabilities. Here, we present the first component of this system: Praxish, a reconstruction of the low-level exclusion logic language atop which the rest of Versu’s functionality is based.

Introduction

Versu (Evans and Short 2014a) was an ambitious authoring framework intended to enable the creation of highly responsive interactive dramas. It was created by the interactive narrative systems equivalent of a supergroup, consisting of Emily Short (the narrative designer behind such notable works of highly responsive interactive fiction as *Galatea*); Richard Evans (the AI lead for *The Sims 3*); and Graham Nelson (the creator of the Inform series of domain-specific languages for interactive fiction authoring). Versu began development in 2010 and was acquired in 2012 by Linden Lab (the makers of *Second Life*), where development continued until the project was cancelled in 2014 (Reed 2021).

Several papers and blog posts documenting the Versu framework were released, and a handful of example games made with Versu (including the Regency-era murder mystery *A Family Supper* (Figure 1) and the late Roman Empire thriller *Blood & Laurels*) were briefly made available on iPad. Today, however, both Versu itself and its example games are effectively lost: the games cannot be played on modern versions of the iPad operating system, and the framework has never been released, with Linden Lab declining to sell it back to its original creators or otherwise allow its continued development (Reed 2021).

Unusually among interactive narrative frameworks, Versu was intended to support both a strongly simulationist ap-

proach to authoring (in which fine-grained computational models of social situations and character motivations were allowed to drive the story in potentially unexpected directions) and a detailed, high-quality textual rendering of the story that emerged from gameplay. Players of Versu games could be allowed to drop into the shoes of any character in the story; different non-player characters could be made to play different narrative roles as the player desired, with the character’s underlying personality driving them to make appropriate-seeming decisions in a wide variety of situations; and characters could ideally even be transplanted from one story to another, all without compromising prose quality. Additionally, Versu’s content language Prompter (Nelson 2014) aimed to make it possible for even inexperienced programmers to write simulationist stories. This unusual set of features, combined with the framework’s apparent success at meeting these goals in a handful of medium-scale text games, causes Versu to stand out as something of a high-water mark for interactive drama to this day.

We believe that both intelligent narrative technologies researchers and interactive narrative designers would benefit from the public availability of a system like Versu: the former from its uniqueness as a point in the design space of AI-supported frameworks, the latter from its utility as a suite of tools for crafting responsive stories. For this reason, we are working to release a strongly Versu-inspired authoring framework as a freely available and open-source piece of software. The reconstruction of the entire framework (which consists of several distinct and individually nontrivial software components) is no small task, and so we are currently taking a piecemeal approach: building a complete reconstruction of Versu one significant component at a time, with the most foundational components prioritized first.

In this paper, we apply the methodology of *rational reconstruction* (Tearse et al. 2012) to the Praxis logic language atop which the rest of the Versu framework is based. Rational reconstruction involves thoughtfully reconstructing a piece of software from its technical description to gain insight into what makes it work, as well as into what about it can be improved. In addition, the output of the reconstruction process is a piece of software which can be used as a foundation for future research and development, both by us and by others—an especially important result in this context due to the unavailability of Versu itself.

*These authors contributed equally.



Figure 1: A screenshot of the Versu example game *A Family Supper*, taken from (Evans and Short 2014b). The story so far is laid out like a play script, and what each character is currently thinking about can be viewed by interacting with the character portraits at the bottom of the screen.

Our recreation of Praxis is called Praxish. Praxish is implemented in dependency-free vanilla JavaScript and can consequently be deployed in a wide variety of contexts, including frontend-only mobile and desktop web browser games, as well as in a standalone desktop executable or on the server side via Node.js. Additionally, Praxish is open source.¹ Like Praxis, Praxish is a relatively low-level language that will eventually serve as the compilation target for a higher-level and more approachable content authoring language akin to Versu’s Prompter; however, it also implements a sufficiently full set of features that recognizably Versu-like storyworlds can be created in only the preliminary version of Praxish that exists today. (Figure 2 shows which Versu components have been recreated in Praxish in some form.)

Research Process

Rational reconstruction begins with a thorough review of existing documentation and materials related to the software being reconstructed. The canonical write-up of Versu is a 2014 article in the journal *IEEE Transactions on Computational Intelligence and AI in Games* (Evans and Short 2014a), while the longest and most detailed explication of Versu’s features can be found in an informally published article that was presented at Imperial College London (Evans and Short 2014b). Beyond these articles, we also reviewed the full breadth of documentation available on the Versu

¹<https://github.com/mkremins/praxish>

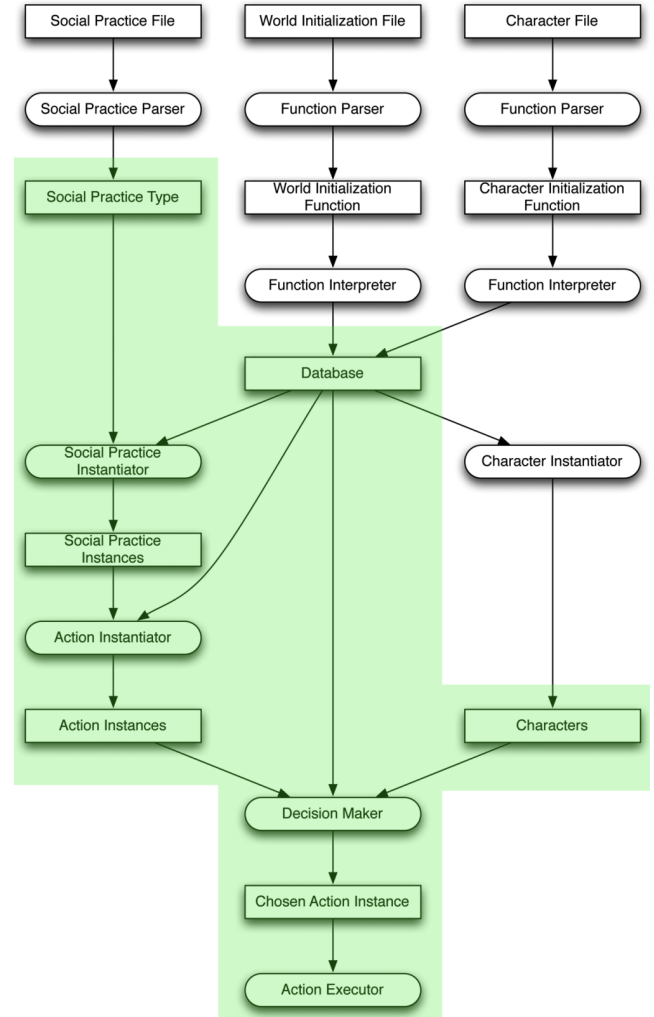


Figure 2: A copy of the Versu system diagram from (Evans and Short 2014a), with a green overlay indicating which components we have reimplemented in Praxish. Our implementations of *characters* and *worlds* are both incomplete: our current example storyworlds manually insert information about characters and the world into the database when run, rather than reading this information from content files. Additionally, we have not yet implemented anything like the higher-level Prompter language for specifying social practice definitions in a more human-readable format, so content authors working with our framework currently have to write practice types in our lower-level JSON format by hand.

website (Short 2014), including a slide deck that gives substantial additional detail on the Praxis language in particular (Evans 2014), and played *Blood & Laurels* on an iPad (belonging to a friend of the Versu developers) that has not been updated since the game was released.

In the process, we found that the most distinctive feature of Versu relative to other narrative-oriented social simulation frameworks (especially at the Praxis layer) is its focus on the separation of content into multiple distinct and composable *social practices*. A social practice provides actions that the various participating agents can choose to perform and tracks the state of a particular social situation, such as a dinner party, a flirtatious side conversation, a game of tic-tac-toe, or a bartender dealing with orders from multiple different customers—all of which could coexist in the same physical space, suggesting appropriate actions to the participating agents in each practice without requiring every practice to be aware of all the others. Practices in Versu were specifically intended to be *role-agnostic* (Evans and Short 2014b)—i.e., to allow any participating character to be played either by an autonomous agent or a human player—and to be potentially reusable across different interactive stories or even entire storyworlds, such that some practices might be equally useful for stories set in realistic, fantasy, and science fictional settings. This grouping of potential character actions into cleanly composable and reusable social practices is therefore a primary goal for our work on Praxish.

We then set about an iterative process of development involving the gradual construction of two example storyworlds and the framework features needed to realize these storyworlds. After reconstructing the most foundational component of the whole framework (the logic database), we split into two teams and attempted to author social practices separately, developing two distinct intermediate JSON formats for representing practice types. We discussed the differences and similarities between these two formats, merged them into a single format, wrote code to support all of the Praxis language features used in these example practices, and then repeated this two-step process of divergent content authoring and convergent framework code authoring for several iterations, eventually ending up with the set of framework features discussed here. Notably, this resulted in two distinct example storyworlds with conceptually similar but somewhat divergent practice definitions; we discuss these a bit more in the “Case Studies” section below.

Overall, we believe that this process of repeated divergence and reflective convergence resulted in a more robust framework—one that accommodates a wider range of use cases than if we had attempted to construct only a single framework and example storyworld from day one.

Implementation

Praxish consists of three major components: the underlying logic database, which handles the storage, querying, and updating of storyworld state; the social practice layer, which handles the specification and maintenance of the rules that govern character behavior in a particular storyworld; and the characters layer, which handles the specification and performance of the characters present in a particular story.

Logic Database

The logic database is Praxish’s most fundamental component. It stores ground sentences of exclusion logic (Evans 2010)² (i.e., sentences that contain no variables), permits the insertion and deletion of sentences, and can be queried via Prolog-style unification (Siekman 1989) of the various sentences in the database with a provided open sentence (i.e., a sentence containing at least one variable).

Ultimately, the majority of higher-level Praxish functionality consists of running queries against the logic database to discover which actions can be performed by a particular agent and updating the database to reflect any changes associated with the action an agent has chosen to perform.

Data Format As in Praxis, a Praxish sentence is any of:

$$X ::= S \mid S.X \mid S!X$$

...where S denotes a symbol (either a capitalized Variable or a lower-case constant). Sentences are stored in the database as trees of constant symbols.

The dot connector $.$ asserts that the tree on the left-hand side has the subtree on the right-hand side. For instance, the sentence `bedroom.furniture.bed` asserts that `bedroom` has one or more subtrees, including `bedroom.furniture`; which itself has one or more subtrees, including `bedroom.furniture.bed`. This sentence could coexist in the database with sentences like `patio`, `bedroom.occupant.trip`, and `bedroom.furniture.desk`, allowing for the specification of a hierarchical world model.

The exclamation point connector $!$ is like the dot connector, but asserts that the left-hand side has only a single subtree, namely that specified by the right-hand side. Inserting first `bedroom.lockState!locked`, followed by `bedroom.lockState!unlocked`, into the database will result in only the second sentence being retained. However, all other subtrees of `bedroom` will remain unaltered: for instance, `bedroom.furniture` is unchanged.

Querying via Unification The database can be queried by unification of its sentences against an open sentence that contains one or more capitalized variables. Suppose we query with the open sentence `Room.occupant.Person` (which contains two variables, `Room` and `Person`) a database that contains the following sentences:

```
bedroom.furniture.bed
bedroom.occupant.trip
bedroom.lockState!unlocked
patio.occupant.grace
```

This yields two sets of variable bindings: `{Room: bedroom, Person: trip}` and `{Room: patio, Person: grace}`.

²Exclusion logic was later rebranded by Evans as *eremic logic* (Evans 2016). We use the older name here for consistency with the Versu documentation.

Updating the Database The database accepts two basic update commands: `insert` and `delete`. Insertion generally adds a single sentence to the database, but may delete existing sentences representing removed subtrees if a subtree is overwritten by an inserted sentence that contains the exclamation point (!) connector. For instance, if the database initially looks like the following:

```
jukebox.color!red
jukebox.playState.song!allStar
jukebox.playState.songPart!beginning
jukebox.playState.nextUp!closingTime
```

...then inserting `jukebox.playState!broken` will cause the database to look like the following, replacing all information about the jukebox's `playState` with the opaque symbol `broken` (without affecting any other stored information):

```
jukebox.color!red
jukebox.playState!broken
```

Deletion, meanwhile, can be used to remove either a single sentence or all sentences that share a specific prefix. For instance, if the database initially looks like the following:

```
practice.debate.ben.pattie
practice.debate.ben.pattie.whoseTurn!pattie
practice.debate.ben.pattie.topic!agents
practice.greet.trip.gonzalo
```

...then deleting `practice.debate.ben.pattie` will cause the database to look like the following:

```
practice.greet.trip.gonzalo
```

...essentially discarding all information related to the ongoing debate between Ben and Pattie, while preserving everything else.

Deletion in exclusion logic is specifically intended to allow for the easy implementation of *lifetimes* for bundled data. For instance, many social practices (including conversations, games, and so on) store temporary information that is needed to organize the behavior of participating agents while the practice (or a particular temporary state of the practice) remains active, but that can safely be cleaned up as soon as the practice concludes. The ability to easily clean up this transitory state with a single `delete` command is key to keeping the database manageable when many practices are active in parallel.

Ultimately, almost all state updates at higher layers of the Praxish stack are eventually implemented in terms of `insert` and `delete` commands issued to the underlying logic database.

Social Practices

The next layer of Praxish provides affordances for authoring structured *social practices*, which guide the behavior of characters (both player and non-player) by providing these characters with a set of actions they can perform. From a technical perspective, practices are similar to finite state machines: each practice instance encapsulates some state that will be cleaned up when the practice instance is removed;

actions offered by a particular practice are typically made available only when the practice is in a particular state; and taking an action offered by a particular practice will typically advance that practice's state in some way.

Like Praxis, Praxish implements a *constitutive* view of social practices. Unlike the *regulative* view, in which social practices shape a person's behavior by *restricting* what actions they can take, practices in the constitutive view shape a participant's action possibilities in an additive way: the *only* meaningful actions that characters can take are those provided by the practices they are participating in (Evans 2016). In other words, the set of actions available to a character in Praxish is exactly the union of the sets of actions provided by each of the practices in which that character is a participant.

Practice Format Practices are defined as JSON objects with several required and several optional keys. At minimum, a practice must have:

- A string **id** that uniquely identifies this practice type
- A list of string **roles**, containing at least one entry, that contains the names of the practice's uniquely identifying instance variables
- A list of **actions** containing at least one action definition, specifying the actions that this practice provides to its participants (format detailed further below)

Optionally, a practice may also have:

- A human-readable string **name** describing the practice, which may contain variable substitutions from the practice's roles
- A list of static **data** sentences, which are inserted into the database under the `practiceData.PracticeID` prefix when the practice definition is registered
- A list of **init** commands to run when an instance of the practice is spawned, generally used to insert additional instance-level data into the database
- A list of zero or more **functions**, each of which is made available to the `call` command when the practice definition is registered

For an example of a basic practice definition that provides a single action, see Figure 3.

Action Conditions Every action definition consists of two major parts: a list of **conditions** that must be simultaneously fulfilled for the action to be available, and a list of **outcomes** that are executed when the action is performed by an agent. If the conditions for an action can be fulfilled in multiple different ways at the same time, then multiple different versions of the action are made available for the acting agent to choose between.

In considering whether a particular action definition should result in any concrete action possibilities for the currently acting character, the action instantiator first establishes bindings for all of the role variables from the practice that provides the action (for instance, `Greeter` and `Greeted` in the `greet` practice outlined in Figure 3) and for the special variable `Actor` (which always indicates the ID of the currently acting character). It then attempts to unify

```

{
  id: "greet",
  name: "[Greeter] is greeting [Greeted]",
  roles: ["Greeter", "Greeted"],
  actions: [{
    name: "[Actor]: Greet [Other]",
    conditions: [
      "eq Actor Greeter",
      "eq Other Greeted"
    ],
    outcomes: [
      "delete practice.greet.Actor.Other"
    ]
  }]
}

```

Figure 3: A basic practice definition in its entirety. This practice has two roles (Greeter and Greeted) and provides a single action, which allows the greeter to greet the greeted and deletes the practice instance that provides it when performed. An instance of this practice between the characters `trip` and `gonzalo` would be identified in the database by the prefix `practice.greet.trip.gonzalo`.

the action's conditions with one another subject to the initial variable bindings, establishing at least one set of internally consistent bindings for the initially unbound variables in the action's conditions. For every set of bindings that can be established, a single instance of the action is made available.

In the current version of Praxish, several different condition types are allowed:

- Basic propositional conditions, which check that a certain sentence exists in the database
- Equality conditions, which check that two arguments (typically variables) are equal
- Negated conditions, which check that another condition does *not* hold

The original Praxis language implemented a slightly wider range of condition types, including disjunctive conditions, conditions that evaluate mathematical expressions such as greater-than, and conditions containing universal quantifiers. This range of conditions was inspired by those available within PDDL (McDermott 2000), and we eventually intend for our framework to support this full range of features, but mathematical expressions and universally quantified conditions are both left as (relatively straightforward) priorities for future work, since we found no immediate need for them in the example storyworlds that we created so far. Disjunctive conditions can already be implemented in terms of multiple different actions with similar but slightly disjoint conditions and identical player-visible names; this adds authoring burden, but because we intend for Praxish to be a compilation target for a higher-level and more approachable authoring language in the future (and because we found little need for disjunction in authoring our example content), we found that no new Praxish-level features are needed to support disjunction at this time.

Action Outcomes An action's outcomes are simply a list of commands to execute in order when the action is performed. Commands include basic `insert` and `delete` commands that update the underlying logic database directly, as well as a `call` command that can be used to call a named function with zero or more parameters. As described in the following section, the `call` command can be used to implement conditional execution of action effects via functions that execute different commands depending on the values of their parameters.

The `insert` command can also be used to spawn new practice instances. Inserting a sentence into the database of the form `practice.PracticeID.RoleBindings` (for instance, `practice.greet.trip.gonzalo` in the case of a simple greeting practice) will cause a new instance of the named practice to be spawned with the provided role bindings, potentially running the corresponding practice's instance initializer (which can perform additional database updates) in the process. In this way, actions in one practice can spawn subpractices that are used to furnish characters with responses or reactions to an action that was just performed: for instance, asking a question to a specific character in a group conversation practice might spawn a subpractice that provides the targeted character with several different ways of answering (or pointedly refusing to answer) the question that was just asked.

Callable Functions Functions behave similarly to actions, but can be invoked via the `call` command at any time (particularly within the effects of an action or the body of another function) to perform a sequence of parametrized commands. The following `call` command, for instance:

```
call setShouldRespond battler beatrice
```

...would invoke the `setShouldRespond` function (perhaps provided by some sort of group conversation practice) with the parameters `battler` and `beatrice`—i.e., the character who is now socially obliged to respond to another character's direct question and the asker to whom a response is owed.

Functions are also capable of conditionally executing effects. This is because they contain not just a single list of conditions and a single list of effects, but an ordered list of **cases**, each of which has its own separate list of conditions that must be satisfied and effects that will be run if these conditions are met. Execution of a function proceeds down the list of cases, trying to find a case whose conditions all hold; as soon as an appropriate case is found, the body of the case is executed with the bindings established by the function's parameters and the case's conditions, and the search process stops. If no matching case is found, the function does not execute any code.

For example, a function intended to check the state of a tic-tac-toe board after every move and update the game's overall status might have several cases:

- One that matches a victory by horizontal row, and marks the game over with the appropriate player as the victor if one is found
- One that matches a victory by vertical column

- One that matches a victory by diagonal (from the top left to the bottom right)
- One that matches a victory by diagonal (from the top right to the bottom left)
- One that matches a tied board, and marks the game over with no player as the winner

This function could then be called in the effects of every action that updates the board state to ensure that endgame conditions are always marked as soon as they are met.

These case-based functions (which were not present in Praxis) are inspired by, and behave similarly to, pattern-matching functions in functional programming languages like Haskell (Wadler 1987). They allow actions to execute conditional effects without the addition of any new conditional-execution features (beyond unification, which is already used to determine which actions are possible) to the core Praxis language.

Characters

In parallel to the social practice layer, an additional set of Praxis language features supports the definition and performance of non-player characters. This is the most preliminary of Praxis's three existing components and the most likely to change in future versions of the language, but nevertheless already implements a significant portion of the character definition functionality provided by Praxis.

Character Format Characters in the present version of Praxis are defined as JSON objects with several allowed keys, all of which are optional except the **id**:

- An **id** string uniquely identifying this character
- A human-readable **name** string, substituted for the **id** in text presented to the player
- A list of zero or more **goals**, each consisting of a list of **conditions** and an associated scalar **utility** score (described further below)
- A **boundToPractice** string indicating the ID of a social practice to which the character is “bound”, if any (described further below)

Utility-Based Action Selection Non-player characters in Praxis behave by default as reactive utility-based agents that perform a single step of lookahead on the available actions (essentially simulating the immediate consequences of each available action) before deciding what action to perform. Agents choose the action with the highest available utility score to perform; if multiple actions are tied for highest score, one of these is chosen uniformly at random.

Utility scores are assigned to actions on the basis of character *goals*. A character goal consists of two components: a set of **conditions** that must all be simultaneously true for the goal to be satisfied, and a scalar **utility** score that the character “earns” when this goal is satisfied. For instance, a romantic comedy character *kaguya* who wants to impress her crush by ordering the same drink as them at the bar might have a goal represented as follows:

```
{utility: 5,
 conditions: [
  "practice.bar.B.customer.kaguya.order!Drink",
  "kaguya.ship.Crush.evaluation.crush",
  "Crush.preferences.favoriteDrink!Drink"]}]
```

...and consequently would derive 5 utility from any action that results in her placing an order for the drink *Drink* at any arbitrary bar *B* where the *Drink* in question is also the favorite drink of a character she views as a *Crush*.

The overall utility of an action is determined by summing together the utility scores of every character goal that the action would cause to be satisfied. Consequently, characters will generally prefer to take actions that satisfy multiple of their goals at once if such actions are available.

Because goal satisfaction is evaluated via unification in the same way as action or function case conditions, a single goal can be satisfied in multiple different ways. When this is the case, the character earns additional utility for each instance of the goal that is satisfied.

Practice-Bound Agents Sometimes it is useful to include agents in a storyworld that are not traditional human-like social actors, but that are still capable of taking a narrow range of autonomous actions. Examples include both drama managers (which should take only metanarrative actions, perhaps intervening at certain key moments to advance the story to its next high-level plot beat or to adjust the motivations of a central character) and diegetic automata that should not participate in the vast majority of social practices (such as a jukebox, which should only take actions that involve playing through a song selected by another agent).

In Praxis, these kinds of behaviors are implemented by *practice-bound agents*: agents that are restricted to taking only actions provided by a specific named practice. A character definition may specify (via the **boundToPractice** key) the string ID of a practice to which this character should be bound, and the character will thereafter be unable to perform actions provided by any practices besides the specified one. This is a novel feature not available in the original Praxis language, and we discuss its implications in greater detail in the “Key Differences from Praxis” section below.

Player Characters Any agent may be substituted by a human player. Whenever an agent played by the player is allowed to act, the player is asked to directly select an action for the agent to perform from the complete list of actions available to that agent.

Key Differences from Praxis

Though Praxis began as a one-to-one reimplementations of the original Praxis language from available documentation, we soon encountered several points at which it made sense to deviate from Praxis slightly in functionality. This section briefly discusses our key deviations from Praxis and the motivation behind each.

Conditionality via Pattern-Matching Functions

The original Praxis language provides an **if** command that can be used in an action's outcomes (or in the body of a

function) to check the value of a conditional, then execute a list of body commands if the conditional evaluates to true. However, because part of our vision for Praxish is that it will eventually serve as the compilation target for a higher-level and more approachable content language, we would prefer to minimize the number of distinct language features that are supported at the Praxish layer and to push features that are less strictly necessary up the stack to higher-level layers of our framework.

Consequently, we deviate from Praxis in providing only a single Praxish-level conditional primitive: the pattern-matching behavior exhibited in action preconditions, extended to the behavior of callable functions. A function that matches on its parameters and executes a different case depending on these parameters' values can be used to implement arbitrarily complicated conditional effects without necessitating the addition of a separate `if` command, and a higher-level authoring tool or language can directly compile conditional effects in action definitions down to the invocation of an equivalent anonymous function if we decide that we want to implement something like `if` later on.

Practice-Bound Agents

A need for autonomous agents that do not participate fully in normal human social practices (including extradiegetic drama managers and diegetic automata) arose multiple times during our authoring of example content. Though these agents can be implemented by adding an extra condition to almost all actions (such that the vast majority of actions in a storyworld check whether the acting agent is a normal character rather than an automaton before allowing the action to proceed), in practice the need to manually insert these conditions substantially increases the authoring burden (Jones 2023) on practice authors. Manual insertion of actor type preconditions also limits portability of practices from one storyworld to the next, since there is no universal way to declare to the system that certain agents should not be treated as normal characters. Therefore, though it adds to the overall complexity of the core Praxish language, we still found the addition of a practice-bound agents feature to be worthwhile due to how much authoring burden it alleviates.

Practice-bound agents should also be excluded from selection as *targets* for normal social interactions, since they don't represent typical embodied characters. For instance, characters usually shouldn't try to flirt with the jukebox or pursue a rivalry with the drama manager.³ We initially believed that this would require us to explicitly exclude practice-bound agents from selection as targets of actions outside their bound practice: a nontrivial problem, given that the targets of an action are not explicitly tracked by a language-level feature in Praxish in the way that the currently acting agent is tracked by the `Actor` variable. Fortunately, however, we discovered that the natural style of practice authoring (in which agents must generally take an explicit action to *enter* a practice before they can participate in it, or else must be inserted into a specific role within the

³That said, we would love to play a game in which it makes sense to do either of these things.

practice when the practice is first spawned) tends to exclude practice-bound agents from appearing as potential targets for other practices' actions as well, even without any additional language-level functionality or practice authoring conventions. Because a practice-bound agent will never *initiate* its own participation in other practices, it will generally not be *subject* to actions by other agents in other practices either.

In a Praxish-based game that permits players to select which character to play as, developers may allow players to select a practice-bound agent as the player character. This enables the creation of unusual social simulation-driven play experiences such as *Juke Joint* (Ryan et al. 2016), in which the player plays as a haunted jukebox with very limited affordances for interaction, and *Cozy Mystery Construction Kit* (Kreminski et al. 2019), in which one of the players is cast in the role of a drama manager who takes exclusively extradiegetic rather than in-character actions.

No Typechecking

Praxis was a statically typed language that enforced type annotations (provided by content authors) on the symbols that could be bound to certain roles, both in social practices and as arguments to callable functions. Praxish, on the other hand, makes no attempt to implement typechecking or type annotations, largely for the sake of language simplicity and ease of implementation. Though we believe that typechecking would eventually be nice to have as an assistance feature for content authors, we plan to implement typechecking later in the higher-level authoring tools rather than include it in the lower-level Praxish language, at least for now.

Case Studies

As discussed in the "Research Process" section above, we split into two content authoring teams during the research process and created two distinct but similar example storyworlds while developing Praxish. Both feature the same general setting (a bar that contains a jukebox and some space for playing tabletop games) and both feature some shared social practices, though they differ in the details.

Each storyworld also presents a slightly different user interface. One, created to test autonomous character behavior, simply simulates and prints out a sequence of autonomous character actions in the browser console to illustrate the nondeterministic and interleaved progression of the various social practices active in the world. The other (Figure 4), created to test player interaction, allows the player to create several named characters, then illustrates simple round-robin turn-taking in which the player is repeatedly allowed to choose the next action of each character.

Together, these example storyworlds (both available as part of this paper's associated software artifact) demonstrate a wide variety of Praxish functionality:

- The `greet` practice demonstrates how practices provide actions only to their participants; how multiple instances of the same practice type can exist simultaneously between different sets of characters; and how a practice can remove itself from the database once it concludes.

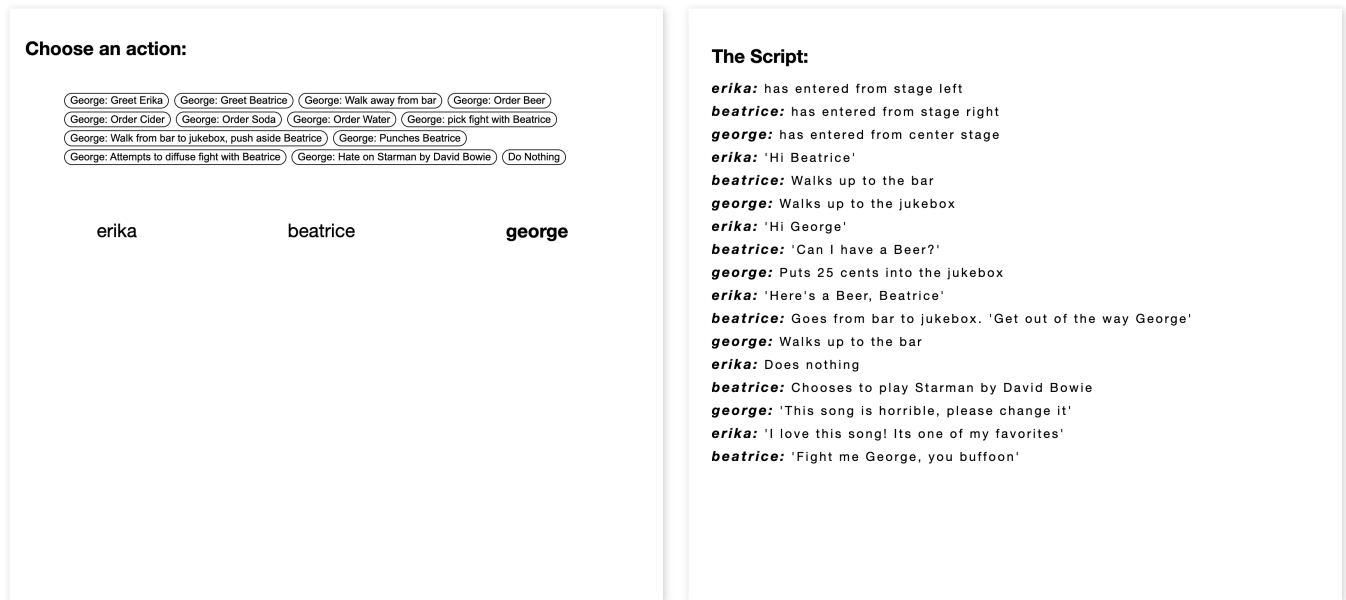


Figure 4: A screenshot of our interactive demo, showing the actions available to the currently acting character on the left and the transcript of the story so far on the right.

- The `tendBar` practice demonstrates role asymmetry within practices (the bartender has access to a very different set of actions than the other characters), as well as the capacity for characters to independently *join* and *leave* practices by walking up to and away from the bar. Two characters in the non-interactive demo also exhibit preferences for specific kinds of drinks, while the bartender has a preference for serving one of the customer characters over the other.
- In the non-interactive demo, the `ticTacToe` practice makes extensive use of a callable pattern-matching function to determine whether an endgame state has been reached every time a character makes a move.
- In the non-interactive demo, the `jukebox`, which plays through a requested song autonomously once another character makes a request, is a practice-bound agent: it can take actions, but only those provided by the `jukebox` practice to which it is bound, and it will never be selected as the target for normal social interactions.
- In the interactive demo, the `fight` practice demonstrates both how practices can be spawned by other practices, as well as how multiple different causes can all lead to the same effect: there are several different ways for characters to become angry at one another, but all of them can result in a bar fight if the angered character decides to escalate to a physical attack.

Both example storyworlds contain only a few hundred lines of practice and character definition JSON, split across four distinct practice types in the non-interactive demo and six in the interactive demo. We found that our pace of content authoring dramatically accelerated as framework features solidified, eventually enabling the development of multiple complex practices per day.

Related Work

We are aware of two prior non-Versu projects that made use of features from the Versu AI architecture: the now-defunct Character Engine by Spirit AI (a company whose founding team included Versu’s Emily Short) made use of Praxis-like social practices to structure agent behaviors, while the experimental AI-based game *MKULTRA* (Horswill 2018) was built around an exclusion logic database. Unfortunately, the Character Engine framework is now unavailable, and *MKULTRA* does not replicate features of the Versu architecture beyond the foundational database layer.

More broadly, there exist a number of logic-based domain-specific languages for specifying the behavior of characters in narrative-oriented social simulations. These all differ substantially from Praxis in terms of compositionality: where Praxis (like Praxis before it) aims to cleanly separate social behaviors into distinct social practices that can be reused from one story to the next, instantiated multiple times within a single story, and potentially even ported from one storyworld to another, other systems typically lump together all of the actions that are possible for characters to take into a single storyworld definition that does not consist of multiple readily separable practice definitions.

Ceptre (Martens 2015) is a logic language for interactive storytelling based on *linear* rather than exclusion logic, which provides a different set of solutions to the problems of modeling stateful social situations. Actions in Ceptre (called *rules*) may produce and consume *resources* that influence which actions are and are not available, and rules are grouped into *stages* that can be used to represent different high-level social contexts or successive states of a progressing social practice. However, the story may not be in multiple different stages at the same time, so stages cannot

be used to cleanly separate social behaviors into composable or reusable practices in the way that a Versu-like approach to practice definitions would permit. Additionally, Ceptre does not provide any built-in mechanism for intelligent decision-making by non-player characters when multiple actions are possible. Ceptre does offer an approachable graphical authoring tool (Card and Martens 2019) that may serve as inspiration for our own future work on higher-level authoring tools for Praxish.

The *Comme il Faut* line of systems (McCoy et al. 2014), including successor systems *Ensemble* (Samuel et al. 2015) and *Kismet* (Summerville and Samuel 2020), are all based on a two-stage model of character decision-making. Characters in these systems first form high-level *volitions* toward one another (generally relationship status intents such as “start dating” or “become rivals”), with their propensity to pursue specific volitions calculated via storyworld-specific *influence rules*, and then take actions that are specified to advance these volitions. These volitions, however, are relatively limited: unlike in Praxish, characters cannot intend arbitrarily complicated sets of logical conditions, limiting the expressive flexibility of character goals. Nevertheless, *Ensemble* in particular is associated with a graphical authoring tool that may further influence our own future work on Praxish-based authoring tools.

Finally, the recently released social simulation game *City of Gangsters* (Zubek et al. 2021) is backed by a logic language called BotL that is largely used for *reactive* querying and state updates. Queries are periodically executed against the state of the world to discover whether certain logical patterns are met, in a fashion similar to story sifting (Kreminski, Wardrip-Fruin, and Mateas 2023), and relationship states for certain characters are updated if they are. Unlike the other systems discussed here (including Praxish), BotL is not responsible for handling the complete character action selection loop in *CoG*, and nothing like the definition of distinct, composable and reusable social practices is attempted. However, *CoG* does provide an interesting alternative take on how social norms could be encoded in a logic language for narrative-oriented simulation.

Conclusion and Future Work

We have presented Praxish: an open-source rational reconstruction of the Praxis logic language, which served as the foundation for the ambitious and now-defunct proprietary interactive narrative authoring framework Versu. Praxish can be used to define interactive storyworlds in terms of multiple distinct, composable, reusable and role-agnostic *social practices* that can operate in parallel to provide characters with a wide variety of action possibilities, as well as highly expressive character *goals* that drive reactive utility-based action selection by non-player characters.

In the future, we plan to continue our Versu reconstruction efforts by introducing a higher-level and more approachable authoring language that compiles down to Praxish, akin to Versu’s Prompter (Nelson 2014). This may be accompanied by other approachable authoring tools, including tools for creating NPCs that go beyond the original authoring affordances of Versu itself.

We also plan to create one or more larger storygames in Praxish as a case study. In particular, we aim to try reproducing a game that is similar in scale and storyworld to either *A Family Supper*, *Blood & Laurels*, or another Versu example game, to achieve and eventually demonstrate feature-parity with the original Versu framework. These efforts may be accompanied by one or more authoring studies (Hargood and Green 2023) intended to evaluate the usability of Praxish and our other authoring tools for content authors.

Acknowledgements

Thanks to Jacob Garbe for preserving, and allowing us to play, a working installation of the Versu example game *Blood & Laurels*.

References

- Card, A.; and Martens, C. 2019. The Ceptre Editor: A Structure Editor for Rule-Based System Simulation. In *2019 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*, 133–137. IEEE.
- Evans, R. 2010. Introducing Exclusion Logic as a Deontic Logic. In *DEON*, volume 10, 179–195. Springer.
- Evans, R. 2014. Praxis: A Logic-Based DSL for Modeling Social Practices. <https://versublog.files.wordpress.com/2014/05/praxis.pdf>.
- Evans, R.; and Short, E. 2014a. Versu—a simulationist storytelling system. *IEEE Transactions on Computational Intelligence and AI in Games*, 6(2): 113–130.
- Evans, R. P. 2016. Computer models of constitutive social practice. *Fundamental Issues of Artificial Intelligence*, 391–411.
- Evans, R. P.; and Short, E. 2014b. The AI Architecture of Versu. <https://versublog.files.wordpress.com/2014/05/versu.pdf>.
- Hargood, C.; and Green, D. 2023. The authoring tool evaluation problem. In *The Authoring Problem: Challenges in Supporting Authoring for Interactive Digital Narratives*, 303–320. Springer.
- Horswill, I. 2018. Postmortem: MKULTRA, an experimental AI-based game. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, volume 14, 45–51.
- Jones, J. D. 2023. Authorial Burden. In *The Authoring Problem: Challenges in Supporting Authoring for Interactive Digital Narratives*, 47–63. Springer.
- Kreminski, M.; Acharya, D.; Junius, N.; Oliver, E.; Comp-ton, K.; Dickinson, M.; Focht, C.; Mason, S.; Mazeika, S.; and Wardrip-Fruin, N. 2019. Cozy Mystery Construction Kit: prototyping toward an AI-assisted collaborative storytelling mystery game. In *Proceedings of the 14th International Conference on the Foundations of Digital Games*.
- Kreminski, M.; Wardrip-Fruin, N.; and Mateas, M. 2023. Authoring for story sifters. In *The Authoring Problem: Challenges in Supporting Authoring for Interactive Digital Narratives*, 207–220. Springer.

- Martens, C. 2015. Ceptre: A language for modeling generative interactive systems. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, volume 11, 51–57.
- McCoy, J.; Treanor, M.; Samuel, B.; Reed, A. A.; Mateas, M.; and Wardrip-Fruin, N. 2014. Social story worlds with Comme il Faut. *IEEE Transactions on Computational Intelligence and AI in Games*, 6(2): 97–112.
- McDermott, D. M. 2000. The 1998 AI Planning Systems Competition. *AI Magazine*, 21(2): 35–35.
- Nelson, G. 2014. Prompter: A Domain-Specific Language for Versu. https://versublog.files.wordpress.com/2014/05/graham_versu.pdf.
- Reed, A. A. 2021. 2013: A Family Supper. <https://if50.substack.com/p/2013-a-family-supper>.
- Ryan, J.; Brothers, T.; Mateas, M.; and Wardrip-Fruin, N. 2016. Juke Joint: characters who are moved by music. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, volume 12, 72–78.
- Samuel, B.; Reed, A. A.; Maddaloni, P.; Mateas, M.; and Wardrip-Fruin, N. 2015. The Ensemble engine: Next-generation social physics. In *Proceedings of the Tenth International Conference on the Foundations of Digital Games (FDG 2015)*, 22–25.
- Short, E. 2014. How Versu works. <https://versu.com/about/how-versu-works/>.
- Siekmann, J. H. 1989. Unification theory. *Journal of Symbolic Computation*, 7(3-4): 207–274.
- Summerville, A.; and Samuel, B. 2020. Kismet: a small social simulation language. In *ICCC-WS 2020: Joint Workshops of the International Conference on Computational Creativity*.
- Tearse, B.; Mawhorter, P.; Mateas, M.; and Wardrip-Fruin, N. 2012. Lessons learned from a rational reconstruction of Minstrel. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 26, 249–255.
- Wadler, P. 1987. Views: A way for pattern matching to cohabit with data abstraction. In *Proceedings of the 14th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages*, 307–313.
- Zubek, R.; Horswill, I.; Robison, E.; and Viglione, M. 2021. Social modeling via logic programming in City of Gangsters. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, volume 17, 220–226.